

Keval Shah

New York, NY | (607) 235-1181 | keval.shah61298@gmail.com | [linkedin.com/in/keval-shah0612](https://www.linkedin.com/in/keval-shah0612) | github.com/kevalshah0612

EDUCATION

Binghamton University, State University of New York | Binghamton, NY

Jan 2025 - May 2026

Master of Science, Computer Science (AI Specialization), GPA: 4.00/4.00

- **Teaching Assistant (Aug 2025- Present):** Guided 120+ students through Java, C++, SQL, and database systems as Teaching Assistant; improved code submission quality by 30% over one semester

Gujarat Technological University | Ahmedabad, India

Aug 2016 – Sep 2020

Bachelor of Engineering, Computer Engineering

PROFESSIONAL EXPERIENCE

Global Health Impact | Software Engineering Intern, New York, NY

Jun 2025 – Aug 2025

- Built Python ingestion pipelines processing 10M+ weekly WHO health records into PostgreSQL and MongoDB, eliminating 90% of manual data preparation for research analysts
- Developed CSV extraction scripts for country, drug, and disease datasets, generating analysis-ready inputs that powered downstream query and dashboard workflows
- Shipped FastAPI REST APIs on AWS for real-time disease burden queries across 150+ countries, cutting reporting turnaround from hours to under 30 seconds
- Launched React research dashboard with world-map exploration and multi-axis filters, replacing manual spreadsheet exports for 6 global research teams

Tata Consultancy Services | Software Engineer II, Gandhinagar, India

Mar 2021 – Dec 2024

- Automated Python and Docker cloud workflows across AWS and GCP, reducing manual infrastructure overhead by 30% across 10 production enterprise applications
- Instrumented Datadog and CloudWatch telemetry into production systems, cutting incident diagnosis time from hours to minutes across 3 connected applications
- Implemented Spring Security RBAC across 10 applications, closing user-access gaps and reducing support tickets by 22%
- Restored a failing Java production service with Redis caching within 2 months, recovering 30% throughput and earning a Wabtec Client Award

PROJECTS

[EvalTrace](#) | Python, DeepEval, LangChain, GitHub Actions, RAG Evaluation

- Built automated LLM regression pipeline with DeepEval that runs on every pull request via GitHub Actions, reducing RAG hallucination rate from 23% to 4%
- Enforced relevance thresholds and answer-quality gates that block low-scoring retrieval changes before merge, making LLM quality measurable in CI/CD

[FilingQuery](#) | Python, LangChain, HuggingFace, pgvector, FastAPI, PostgreSQL

- Architected RAG system over 5,000+ SEC filings using LangChain and pgvector, answering financial questions with cited source passages for auditability
- Reduced financial document lookup from hours to under 30 seconds by combining semantic embedding search with keyword retrieval over a structured PostgreSQL corpus

[ReviewBot](#) | Python, FastAPI, LangGraph, Redis, GitHub Actions

- Designed LangGraph multi-agent workflow with Redis memory that preserves full repository and PR context across sequential code analysis and review steps
- Deployed FastAPI endpoints integrated with GitHub Actions CI, returning structured style, security, and test feedback on every pull request before merge

TECHNICAL SKILLS

AI and ML: Python, machine learning, scikit-learn, RAG, embeddings, vector search, model evaluation

AI Frameworks and APIs: LangChain, LangGraph, HuggingFace, DeepEval, OpenAI API, FastAPI, GitHub Actions

Data Systems: PostgreSQL, MongoDB, Redis, pgvector, Kafka, SQL, unstructured documents

Cloud and Delivery: AWS, Azure, Docker, Linux, CI/CD, Datadog, regression testing